Unsupervised and Unstructured Machine Learning

BA820 – Mohannad Elhamod

BOSTON UNIVERSITY

Boston University Questrom School of Business

Midterm Review

BOSTON UNIVERSITY

Boston University Questrom School of Business

Start
Stop
Continue

BOSTON UNIVERSITY

Please fill this out

- https://forms.gle/VuHJUjtzLrrr1TNY8

Text Mining

BOSTON UNIVERSITY

Boston University Questrom School of Business

Most Data is No Longer Structured...

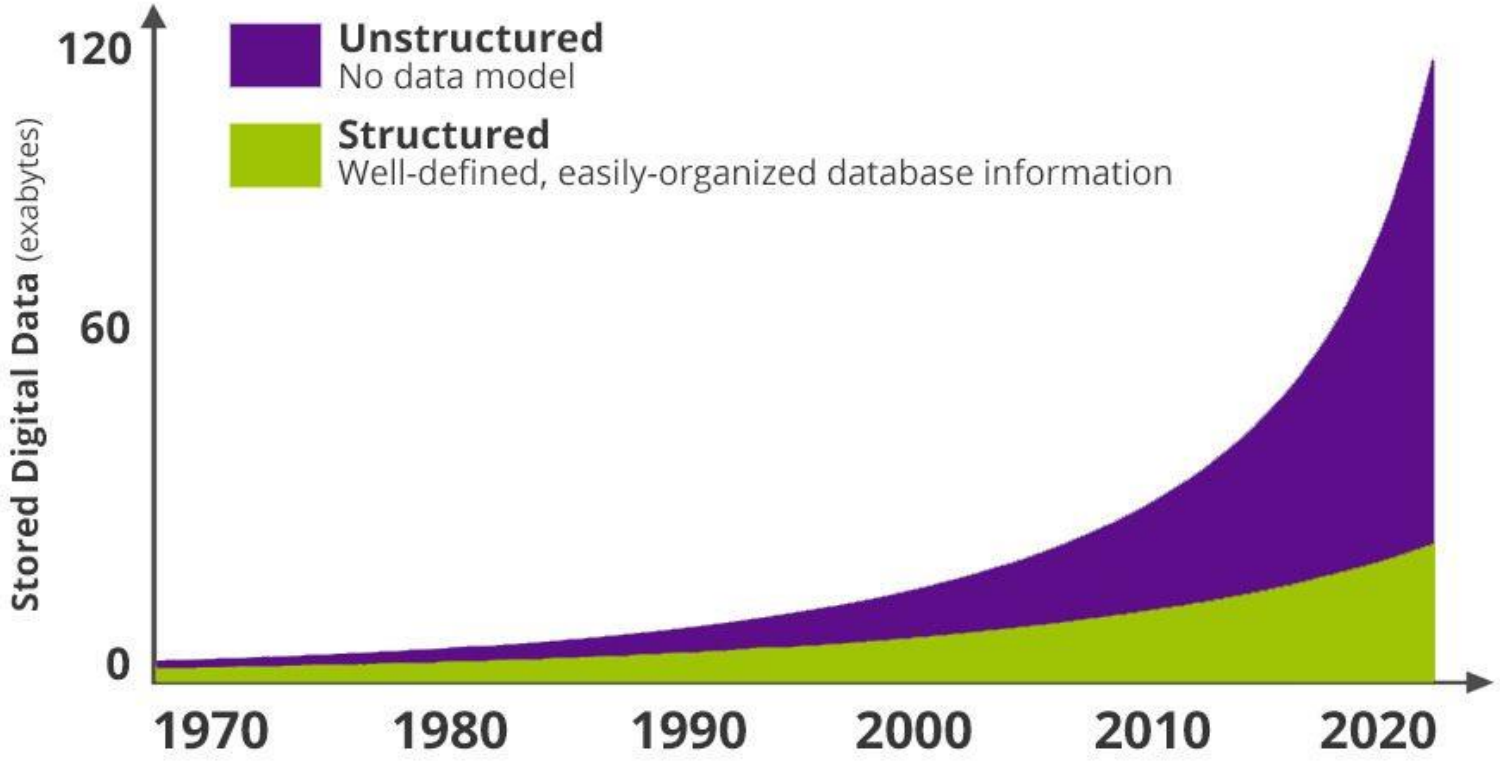

https://www.komprise.com/

What can we do with this data?

- Text classification
- Text generation
- Text summarization
- Music recommendation
- Image categorization
- ….

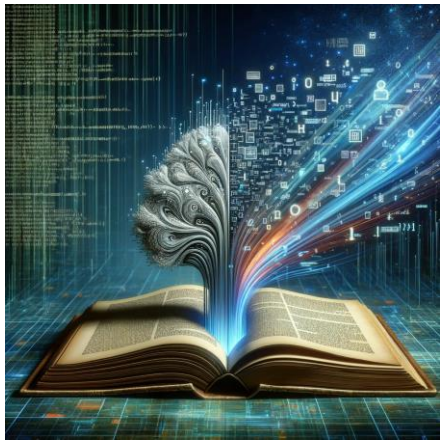

Text Mining

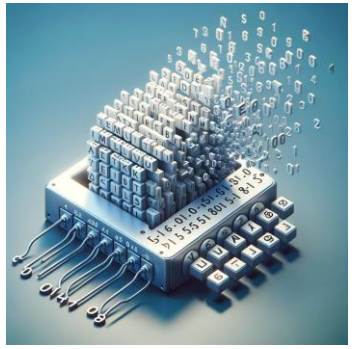

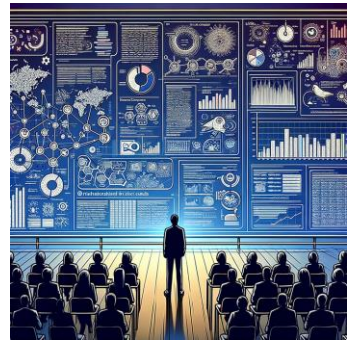

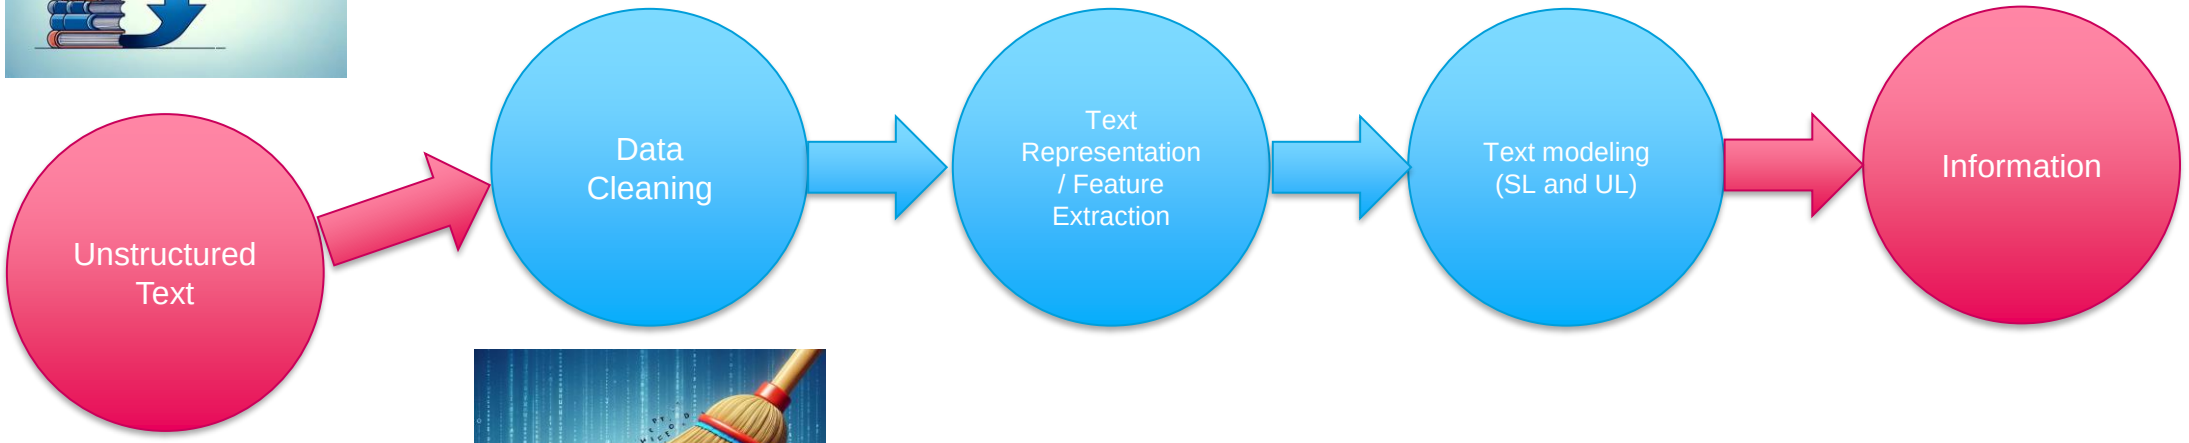
Unstructured Text
Data Cleaning
Text Representation / Feature Extraction
Text modeling (SL and UL)
Information

Collecting Unstructured Data

- How can we extract this data?
  - Datasets found publicly (UCI, Kaggle, Common Crawl, Wikimedia Downloads, etc.).
  - Using APIs (e.g., twitter API, News API).
  - Web scrapping (e.g., BeautifulSoup, Selenium)
  - Own private data.

Data Cleaning with *Regex*

- Just like structured/tabular data, we generally need to clean up the text to make it more useful.
  - Examples: case-sensitivity, punctuation, etc.

- *Regex* is used for finding string matches and formatting text:
  - `Playground:` https://www.w3schools.com/python/python_regex.asp
  - `Cheat sheet:` https://www.debuggex.com/cheatsheet/regex/python

How to Represent Text?

- Computers do not understand text…
- We need to represent text in a language they understand… Numbers!

- Simple proposals (we are just brainstorming):
  - Each sentence is represented in terms of the words it contains…
    - This is called ***Tokenization.***
  - Each word is represented by a number…
    - This is called ***Vectorization***.

Tokenization: Some Terminology

- ***Document:*** A body of text (e.g., a tweet, a pdf, an article, etc.).
- ***A token:*** The building block of a document.
  - Examples: character-level, word-level, …
- ***A separator:*** Special tokens that split a document into tokens.
  - Examples: punctuation, spaces,
- Demo

Need for Advanced Text Pre-Processing

- Simply extracting tokens does not preserve meaning/semantics.
- Some words occur too frequently in any text. These are called *stop words* and are generally removed.
- Issues:
  - ***Stemming:*** big, bigger, biggest
  - ***Lemmatization:*** drive drove driven
  - ***Homonyms:*** bank (river or money?)
  - ***Synonym:*** Yes, sure.
- We will come back to this later…

Text Modeling

- Once text data is in the proper representation (i.e., tokenized and vectorized), we can apply the methods we have learned so far:
  - Unsupervised ML (e.g., dimensionality reduction, clustering, etc.).
  - Supervised ML (e.g., classification, translation, etc.).
- **Demo**

Text Vectorization

BOSTON UNIVERSITY

Vectorization: Text as Numbers

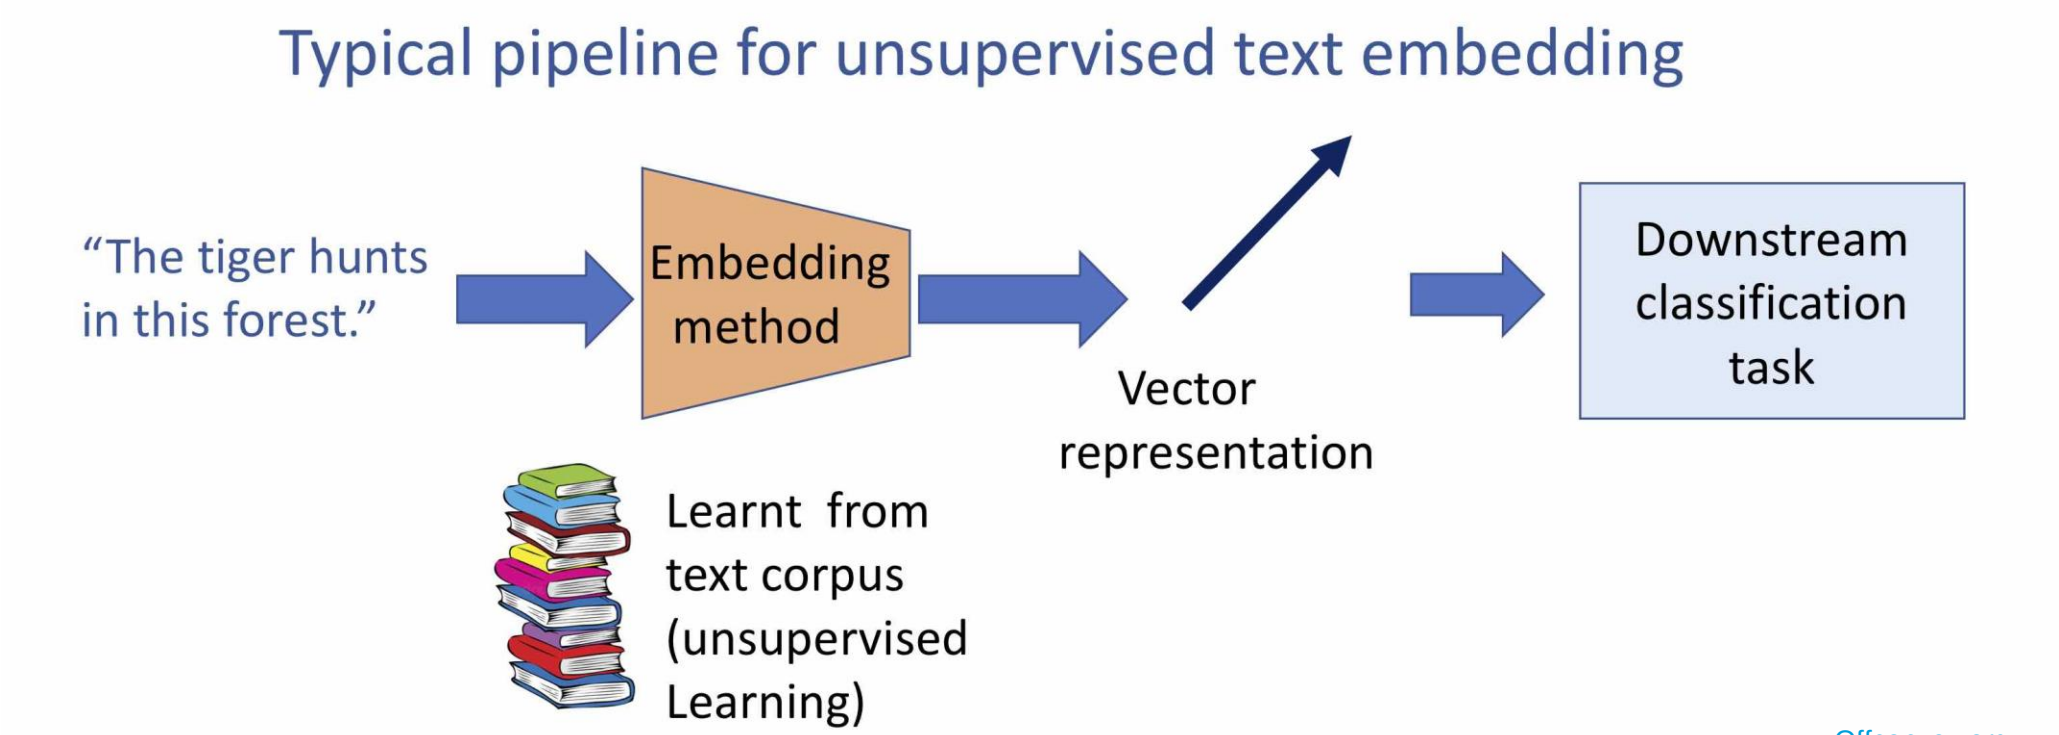

Offconvex.org

How Would You Represent a Document?

- Let's start simple. Represent a document simply as *a collection* of *tokens*.
  - Vectorization by document (not token).
- This approach is called ***Bag of Words (BoW).***
- **Demo**

Sentence	hockey	fun	i	like	golf
I like golf!			1	1	1
I like hockey.	1		1	1	
Hockey and golf are fun!	1	1			1

Bag of Words (BoW)

- Cons:
  - Disregards word order.
  - Number of features can be exhaustive.
  - Frequency bias.
    - (e.g., If the word “space” appears in a children’s book, it carries more significance than when it appears in an article about galaxies.
- Pros:
  - Simple to implement

	I	hate	love	golf	soccer
I hate golf and love soccer	1	1	1	1	1
I hate soccer and love golf	1	1	1	1	1

Document Similarity

- How do we measure if two documents are similar?
  - We need a metric like Euclidean distance.
  - But… What if documents have different lengths?

- We need a metric that is robust to differences in document size…
  - Enter *Cosine Similarity.*

Offconvex.org

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\|\|\mathbf{B}\|} = \frac{\sum_{i=1}^{n} A_i B_i}{\sqrt{\sum_{i=1}^{n} A_i^2}\sqrt{\sum_{i=1}^{n} B_i^2}},$$

Document Similarity

Euclidean distance vs Cosine similarity

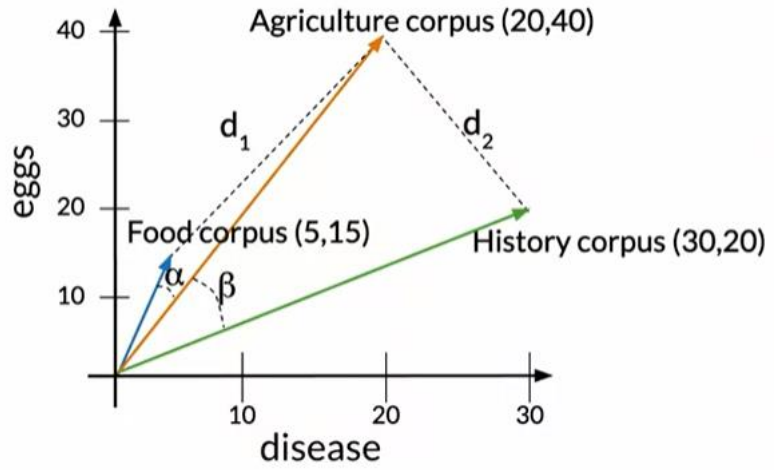

Euclidean distance: $d_2 < d_1$

Angles comparison: $\beta > \alpha$

The cosine of the angle between the vectors

$$\text{similarity} = \cos(\theta) = \frac{\mathbf{A} \cdot \mathbf{B}}{\|\mathbf{A}\|\|\mathbf{B}\|} = \frac{\sum_{i=1}^{n} A_i B_i}{\sqrt{\sum_{i=1}^{n} A_i^2}\sqrt{\sum_{i=1}^{n} B_i^2}},$$

medium.com